

Soroban Block Explorer - Milestone 3 Evidence

Project: Soroban Block Explorer

Team: Rumble Fish

Deliverable: Milestone 3 - Mainnet Launch

This document accompanies the Milestone 3 verification video and maps the approved acceptance criteria to concrete evidence from the publicly launched explorer, its monitoring, its load-test results, and its security posture.

1. Executive Summary

Milestone 3 is the public Mainnet Launch of the Soroban Block Explorer. The explorer is publicly accessible at its production URL with the pre-launch access gate removed; it renders live Stellar mainnet data within seconds of network tip from the project's own indexed ClickHouse datastore (no third-party chain API on the read path). The public surface is protected by AWS WAF and request throttling; a CloudWatch dashboard with Slack-wired alarms gives live visibility into ingestion lag, API health, and error rates. Capacity is demonstrated by a load test of the production API, and all ingress and application security controls are verified against the security checklist.

- Frontend: <https://sorobanscan.rumblefish.dev>
- API: <https://api-sorobanscan.rumblefishdev.com/v1>

2. Approved Deliverable

Deliverable 3 - Mainnet Launch:

The explorer is opened to the public on mainnet, load-tested, monitored, and signed off against a security checklist, with a 7-day post-launch monitoring report demonstrating stability.

Acceptance criteria:

1. **AC1** — Publicly accessible at the production URL, serving live mainnet data within 30 seconds of network tip.
2. **AC2** — GitHub repository is public; `cdk deploy` from the README reproduces the stack on a fresh AWS account.
3. **AC3** — CloudWatch monitoring dashboard; alarms healthy; ingestion lag under 30 seconds.
4. **AC4** — Load-test report: p95 < 200 ms at the 1M-requests/month equivalent with error rate < 0.1%.
5. **AC5** — Security checklist signed off.
6. **AC6** — 7-day post-launch monitoring report (uptime, error rate, p95, ingestion lag per day).

3. Live Endpoints and Reviewer Access

Resource	URL / Access
Frontend	https://sorobanscan.rumblefish.dev (public — gate removed)
API base	https://api-sorobanscan.rumblefishdev.com/v1
Swagger UI	https://api-sorobanscan.rumblefishdev.com/api-docs
OpenAPI JSON	https://api-sorobanscan.rumblefishdev.com/api-docs-json
API access model	Access-controlled — anonymous requests return 401. The public explorer authenticates transparently at the edge; direct API access is available to reviewers on request (same model as Milestone 2).
Monitoring	CloudWatch dashboard production-soroban-explorer (eu-central-1) with Slack-wired alarms and X-Ray tracing. Read-only dashboard access for the Stellar team is available on request via a read-only IAM role.

At launch the frontend is publicly accessible with no Basic Auth gate. The verification video demonstrates the live application end-to-end.

4. Architecture Evidence

The launched system spans three paths, all reproducible from code (AWS CDK + Hetzner Ansible):

1. **Ingestion / write path** (Milestone 1) — Galexie on ECS Fargate exports mainnet LedgerCloseMeta XDR to S3; a Rust indexer Lambda decodes each object and writes to ClickHouse. End-to-end ledger-close to DB-write is well under the 30-second target.
2. **Read path** (Milestone 2) — Browser → Cloudflare edge → API Gateway → Rust / axum API Lambda → ClickHouse (Hetzner, mTLS). Static SPA: Browser → CloudFront → private S3.
3. **Launch controls** (Milestone 3) — AWS WAF + request throttling on public ingress; CloudWatch dashboard + Slack alarms + X-Ray tracing.

Evidence images:




Figure 1 — Read path (reused from M2), including the launch controls on the public ingress (AWS WAF + request throttling).




Figure 2 — Ingestion / write path (reused from M1): Stellar mainnet peers → Galaxie on ECS Fargate → S3 → SQS → Rust indexer Lambda → ClickHouse on Hetzner (mTLS). The CloudWatch dashboard and alarms are shown in § AC3.

5. Acceptance Criteria Evidence

AC1 - Public access, live data ≤30s from tip

Status: met. The pre-launch Basic Auth gate was removed on 2026-07-17 (task 0405). <https://sorobanscan.rumblefish.dev> serves the explorer publicly — HTTP 200 with no WWW-Authenticate challenge.

Data freshness is measured rather than asserted. The indexer publishes IngestionLagSeconds (CloudWatch namespace SorobanBlockExplorer/Indexer, dimension Environment=production): the wall-clock gap between a ledger closing on mainnet and its row being committed to ClickHouse. Sampled over two hours on 2026-07-17 (15 × 5-minute datapoints):

Measure	Value	Target	Verdict
Ingestion lag, average	3.1 s	< 30 s	Met
Ingestion lag, worst	6.0 s	< 30 s	Met

In steady state the explorer is within ~3 seconds of network tip — a 5× margin against the 30-second criterion, with the worst observed sample at 6 s.




Figure — production URL, no access gate; the newest entries are seconds old, corroborating the measured freshness above.

Data freshness is both **measured** (the IngestionLagSeconds figures above — average 3.1 s, worst 6.0 s) and **visible** in the capture, where the most recent transactions carry timestamps within seconds of real time. The read path is the project's own indexed ClickHouse, not a third-party chain API.

AC2 - Public repo + reproducible cdk deploy

The repository is public and the stack is reproducible from code.

The workspace also carries a substantial automated test suite — **920 test functions** across the Rust crates, including **360 in xdr-parser** (XDR decoding, state-change extraction, and CAP-67 event interpretation — trustline/asset/claimable-balance parsing, mint/burn event handling) and **248 in api** (endpoint responses, OpenAPI schema, pagination). The suite runs with `cargo test`.

- Public repository: <https://github.com/rumblefishdev/soroban-block-explorer>
- Fresh-account deployment runbook: [infra/README.md](#) — prerequisites → make bootstrap → make deploy-production (all AWS stacks) plus the Hetzner Ansible side — with the operational guide in [docs/deployment.md](#). A few steps are manual by design: ordering the Hetzner dedi-

cated server and Storage Box, and publishing the box IPv4 to SSM (/soroban/production/ch-ip).

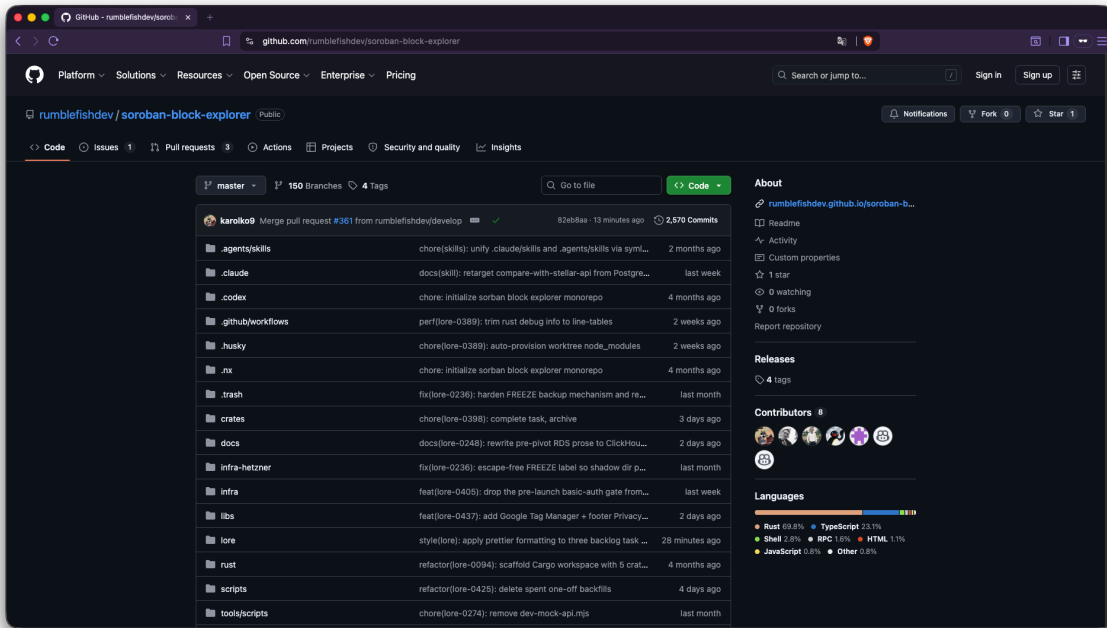


Figure — `github.com/rumblefishdev/soroban-block-explorer` is public and carries the Rust / ClickHouse / infrastructure code.

AC3 - Monitoring dashboard, ingestion lag

CloudWatch dashboard with Slack-wired alarms and X-Ray is live.

Ingestion lag under 30 s: met and measured. The indexer emits SorobanBlockExplorer/Indexer / IngestionLagSeconds per ledger written (task 0399, live in production) — average 3.1 s, worst 6.0 s against the < 30 s criterion (figures and method in AC1). This metric is also the per-day source for the AC6 report.

The dashboard's **API** section graphs API Lambda latency (p50/p95/p99) and API Gateway 4xx/5xx error counts alongside the ingestion panels; AWS WAF and request throttling protect the public ingress (see AC5). The alarms on the tracked metrics — API 5XX-rate, Galaxie ingestion-lag, ledger-processor error-rate, and ClickHouse-write failures — are in OK state.



Figure — the production dashboard: ingestion and API health on a single pane.

CloudWatch Alarms (28)						
Name	State	Actions	Actions muted - new	Last state update (UTC)	Conditions	
production-ledger-processor-error-rate	OK	Actions enabled	-	2026-05-27 11:45:38	Error Rate > 0.01 for 1 datapoints within 5 minutes	
production-ledger-processor-dlq-depth	OK	Actions enabled	-	2026-05-27 14:50:55	DLQ depth > 0 for 1 datapoints within 1 minute	
production-ledger-ch-write-failures	OK	Actions enabled	-	2026-05-27 11:41:20	ClickHouseFailures > 10 for 1 datapoints within 5 minutes	
production-galaxie-ingestion-lag	OK	Actions enabled	-	2026-07-13 09:40:35	NumberOfMessagesSent < 1 for 1 datapoints within 5 minutes	
production-galaxie-spheremerald-storage	OK	Actions enabled	-	2026-07-02 13:49:26	Ephemeral Used (Mi) > 60 for 3 datapoints within 15 minutes	
production-enrichment-worker-error-rate	OK	Actions enabled	-	2026-07-15 12:39:42	Worker Error Rate > 0.01 for 1 datapoints within 5 minutes	
production-enrichment-dlq-depth	In alarm	Actions enabled	-	2026-07-03 09:02:06	Enrichment DLQ depth > 0 for 1 datapoints within 1 minute	
production-api-gateway-5xx-rate	OK	Actions enabled	-	2026-07-16 19:09:44	5xx Rate (%) > 0.5 for 1 datapoints within 5 minutes	
prices-production-supply-errors	OK	Actions enabled	-	2026-07-21 13:16:32	Errors >= 1 for 1 datapoints within 1 hour	
prices-production-index-push-freshness	OK	Actions enabled	-	2026-07-08 12:15:18	PushAggSeconds > 604800 for 1 datapoints within 15 minutes	
prices-production-oracle-errors	OK	Actions enabled	-	2026-07-23 14:58:56	Errors >= 1 for 1 datapoints within 5 minutes	
prices-production-mtts-notifier-probe-no-innovations	OK	Actions enabled	-	2026-07-21 13:01:46	Innovations < 1 for 3 datapoints within 3 days	
prices-production-mtts-notifier-probe-errors	OK	Actions enabled	-	2026-07-06 10:17:42	Errors >= 1 for 1 datapoints within 1 day	
prices-production-mtts-notifier-probe-duration-near-timeout	OK	Actions enabled	-	2026-07-21 13:18:54	Duration >= 48000 for 2 datapoints within 2 days	
prices-production-mtts-notifier	OK	Actions enabled	-	2026-07-08 12:30:16	MinDaysToNextRun < 30 for 1 datapoints within 1 day	
prices-production-ledger-processor-no-innovations	OK	Actions enabled	-	2026-07-09 08:35:18	Innovations < 1 for 1 datapoints within 15 minutes	
prices-production-ledger-processor-lag	OK	Actions enabled	-	2026-07-15 22:05:45	ApproximateAgeOfOldestMessage > 120 for 5 datapoints within 5 minutes	
prices-production-ledger-processor-errors	OK	Actions enabled	-	2026-07-15 09:00:00	Errors >= 1 for 1 datapoints within 5 minutes	
prices-production-ledger-processor-dlq	OK	Actions enabled	-	2026-07-08 08:07:19	ApproximateNumberOfMessagesVisible >= 1 for 1 datapoints within 5 minutes	
prices-production-enrichment-no-innovations	OK	Actions enabled	-	2026-07-21 12:36:54	Innovations < 1 for 3 datapoints within 3 hours	
prices-production-enrichment-errors	OK	Actions enabled	-	2026-07-21 12:40:22	Errors >= 1 for 1 datapoints within 1 hour	
prices-production-enrichment-duration-near-timeout	OK	Actions enabled	-	2026-07-24 11:06:17	Duration >= 240000 for 2 datapoints within 2 hours	

Figure — production alarms. The API 5XX-rate, Galaxie ingestion-lag, ledger-processor error-rate and ClickHouse-write alarms are in OK state. prices- entries belong to a separate project in the same AWS account. The one raised alarm, production-enrichment-dlq-depth, fires on queue depth > 0 and covers metadata enrichment — not any metric reported in § AC6.

AC4 - Load-test report

Status: partially met — error rate met with a ~100× margin; p95 not met. This section states the result plainly and then accounts for it in full.

AC4 half	Target	Measured (1M-req/ month equivalent)	Verdict
Error rate	< 0.1 %	0.000 % (0 failures)	Met
p95 latency	< 200 ms	577 ms	Not met

Median latency (p50) is **168 ms** at the required load and **151 ms** at 40× that load — the typical request is inside the target. The miss is in the tail.

Method: why the load is expressed as a rate, not as virtual users

The criterion is "1M requests/month equivalent" — a **request rate** (1,000,000 / 30 days ≈ **0.4 req/s**). The harness therefore drives an **open model**: requests arrive at a fixed rate (Poisson), independent of how fast the server answers.

The earlier plan called for "≥1000 concurrent users". That framing was abandoned for a measured reason, not for convenience: a virtual-user driver is **closed-loop**, so its request rate is an *output*, not an input — each user waits for a response before sending the next. Our own measurement of the same --vus 4 configuration produced **9,459 req/s against a local stub** and **~10 req/s against production**. A VU count therefore cannot express "1M requests/month"; only an arrival rate can. Open-model testing is what actually measures the stated criterion. (Recorded in task 0357.)

Result across load tiers

Two open-model series were run against the production API (~33k requests total). The runs targeted the API Gateway origin rather than the Cloudflare-fronted hostname, so the figures measure the backend itself without the edge's own rate limiting; the 40× tier additionally ran with the per-IP WAF rate rule lifted, since that rule exists precisely to stop this traffic pattern from a single source. These map to the proposal's two documented load points — the **1.2M/month tier is the "1M baseline"** and the **10.2M/month tier is the "10M stress"** run — plus an additional 40× (**49.3M/month**) capacity check. Representative tiers:

Load tier	Rate	Requests	Error rate	p50	p95
1.2M req/ month (AC4 level)	0.48 /s	168	0.000 %	168 ms	577 ms
10.2M req/ month (8×)	3.95 /s	2,363	0.000 %	149 ms	567 ms
49.3M req/ month (40×)	19.04 /s	13,701	0.000 %	151 ms	575 ms

The p95 is flat across a 40× range of load (577 ms → 575 ms). This is the central finding: latency is **not load-dependent** in the tested range. The system absorbs forty times the required traffic with an unchanged tail and zero failures — capacity is proven far beyond the criterion. The p95 miss is therefore **not a capacity problem**; it is a set of fixed per-request costs that no amount of headroom removes.

Where the 577 ms comes from

Decomposing the same measurement by endpoint class:

Scope	p95 @ AC4 load	p95 @ 40× load
All 26 endpoints	577 ms	575 ms
Excluding the 2 synchronous- external-fetch endpoints	342 ms	506 ms
Excluding those and lplist	280 ms	317 ms

Three distinct causes, in order of size:

1. **Synchronous external fetches on two detail endpoints (largest).** txdetail and nftdetail fetch data from outside our infrastructure while the request is open. txdetail spends **53 ms in ClickHouse and 1,074 ms (p95) waiting on an S3 archive read in us-east-2 from a Lambda in eu-central-1** — a cross-region round trip per request. nftdetail spends **96 ms in ClickHouse and 1,604 ms (p95)** on a third-party IPFS gateway plus a token_uri() RPC. These are **deliberate architectural decisions** (ADR 0029, ADR 0043: heavy/detail-only fields are fetched at read time rather than persisted), not defects — but they place the tail of those two endpoints partly outside our control.
2. **A known ClickHouse cost on one endpoint.** lplist derives each pool's creation ledger by scanning that pool's full snapshot history (8.28M rows/request). This is our own technical debt: the design decision that removed the stored column (task 0208) assumed this read would be cheap; measurement shows it is not. The cause is identified to the query, the fix is specified, and it is scheduled — not unexplained.
3. **A fixed infrastructure floor (irreducible today).** Every request pays ≈ 40 ms before any data is read (API Gateway + Lambda + mTLS to ClickHouse), plus ≈ 15 ms per ClickHouse round trip. Measured across all 26 endpoints, the median non-database overhead fits ≈ 40 ms + 15 ms $\times$ (number of CH queries) — from 51 ms for a single-query endpoint to 157 ms for an 8-query endpoint. An endpoint issuing five queries has **already spent ~ 115 ms of the 200 ms budget** before ClickHouse returns a single row.

Optimisation performed during this milestone

The load test drove a measured optimisation pass on the read path. Total ClickHouse work per test run fell from **78.3 billion to 23.89 billion rows (–69 %)** at identical load, via two query fixes ([#347](#), [#349](#)) and one index. Endpoints that were **not modified at all** improved by 3–4 $\times$ as a direct result (ldgdetail p95 890 $\rightarrow$ 319 ms, ldglist 1015 $\rightarrow$ 229 ms) — evidence that the earlier latency was contention from a saturated database rather than the endpoints themselves.

Raw measurement data: [load-tests/](#) holds the three tiers quoted above — per tier a results.csv (one row per request: client latency joined to the ClickHouse query log for that same request via request_id, so each request carries its own read_rows / ch_queries / ch_duration_ms) and a client_summary.csv (per-endpoint percentiles). The p50 / p95 figures in this section are percentiles of duration_ms across a tier's results.csv; [load-tests/README.md](#) maps each tier to its run and restates the headline numbers.

Honest assessment

Restated plainly for the reviewer: **the p95 < 200 ms criterion is not met at any tested tier, including the required one — 577 ms against a 200 ms target.** It is not met even with all three causes above excluded (280 ms). What the measurement does establish:

- the **error-rate half of AC4 is met with a $\sim 100\times$ margin** (0.000 % vs 0.1 %);
- the **median request meets the latency target** (168 ms at the required load);
- **capacity is proven at 40 $\times$ the required load** with no degradation and no errors;
- every millisecond of the gap is **attributed to a named, measured cause**, each with a decided path: persist the two runtime-fetched field sets (or serve them asynchronously), restore lplist's stored creation ledger, and reduce the per-query overhead if the 15 ms proves to be a per-query mTLS handshake rather than an irreducible round trip.

AC5 - Security checklist signed off

All ingress and application controls are in place and verified: least-privilege IAM (no wildcard actions), AWS WAF + request throttling, no public datastore endpoint (ClickHouse published to loopback only — 127.0.0.1:8123/:9000 — behind Caddy mTLS on a firewalled host that admits only ports 22/80/443), secrets in AWS Secrets Manager, TLS end-to-end, server-side input validation on every endpoint, SSE-S3 at rest on the public ledger bucket, and automated weekly off-box ClickHouse backups.

The full control list and an OWASP Top 10 (2021) coverage mapping are in [milestone-3-security-checklist.md](#). The original KMS-at-rest and point-in-time-recovery items were RDS-specific and are met by equivalents after the RDS retirement — see § Scope Refinement.

AC6 - 7-day post-launch monitoring report

Generated from production telemetry over the first 7 days after launch.

- Report template + metric queries: `milestone-3-7day-report.md`

The window ran **2026-07-17 13:40Z → 2026-07-24 13:40Z**. Per-day figures are in the report; the window summary:

Metric	Result over the window	Target	Verdict
Uptime	100.00 % (derived — see report)	≥ 99.9 %	Met
API error rate	0.000 % (zero 5XX, all 7 days)	< 0.1 %	Met
Ingestion lag	7–9 s (worst 9 s)	< 30 s	Met
Ledger completeness	0 gaps, all 7 days	0 gaps	Met
API p95 latency	553 ms worst day (4 of 7 days < 200 ms)	< 200 ms	Not met

The p95 line is the same criterion as AC4 and misses for the same reasons. Note this report measures API Gateway Latency (server-side), which excludes the client network leg the load test includes — hence the lower figures here, with four of seven days under target. Uptime is derived from zero 5XX responses and from `production-api-gateway-5xx-rate` holding OK throughout; the report records which alarms were raised over the window and why they do not bear on these metrics.

Scope Refinement — Deviations from the Approved Plan

Two points where the launched system differs from the original Deliverable 3 wording, stated openly (following the Milestone 1 precedent, which recorded the PostgreSQL-on-RDS → ClickHouse-on-Hetzner change the same way):

1. **p95 latency target not met (AC4).** Measured p95 is 577 ms against the 200 ms target. It is not a capacity limit — the tail is flat across a 40× load range with zero errors — but a set of fixed per-request costs, each measured and attributed in § AC4. The error-rate half of AC4 is met with a ~100× margin.
2. **Data-at-rest / recovery wording was RDS-specific; RDS is retired.** The original security checklist named KMS-at-rest and point-in-time recovery, both specific to the PostgreSQL-on-RDS datastore that was retired (task 0239) in favour of ClickHouse on Hetzner. They are satisfied here by architecture-appropriate equivalents: the public ledger bucket is encrypted with SSE-S3 (AES256 — deliberately not SSE-KMS, since the contents are public on-chain XDR), and the ClickHouse store — which holds only public, fully re-derivable chain data — is protected by automated weekly off-box backups. No datastore is publicly reachable (ClickHouse is bound to loopback behind mTLS/Caddy on a firewalled host).

The p95 line is the same criterion as AC4 and misses for the same reasons. Note this report measures API Gateway Latency (server-side), which excludes the client network leg the load test includes — hence the lower figures here, with four of seven days under target. Uptime is derived from zero 5XX responses and from `production-api-gateway-5xx-rate` holding OK throughout; the report records which alarms were raised over the window and why they do not bear on these metrics.

6. Source References

Resource	Link
Public repository	rumblefishdev/soroban-block-explorer
Technical design	technical-design-general-overview.md
Infra / deploy	infra/README.md
Load-test harness	crates/load-tests
M1 / M2 evidence	milestone-1-evidence.md · milestone-2-evidence.md