

Soroban Block Explorer - Milestone 3 Evidence

Project: Soroban Block Explorer

Team: Rumble Fish

Deliverable: Milestone 3 - Mainnet Launch

This document accompanies the Milestone 3 verification video and maps the approved acceptance criteria to concrete evidence from the publicly launched explorer, its monitoring, its load-test results, and its security posture.

1. Executive Summary

Milestone 3 is the public Mainnet Launch of the Soroban Block Explorer. The explorer is publicly accessible at its production URL with the pre-launch access gate removed; it renders live Stellar mainnet data within seconds of network tip from the project's own indexed ClickHouse datastore (no third-party chain API on the read path). The data API is fronted by the Cloudflare edge (managed rule sets, rate limiting, bot challenge) with API Gateway request throttling behind it; a CloudWatch dashboard with Slack-wired alarms gives live visibility into ingestion lag, API health, and error rates. Capacity is demonstrated by a load test of the production API, and all ingress and application security controls are verified against the security checklist. The full API reference is published and publicly readable (Swagger UI and OpenAPI JSON, linked in § 3).

Where we stand against the criteria. Five of the six are met. AC4 is met on error rate with roughly a 100× margin (0.000 % against 0.1 %) and **missed on latency — p95 is 577 ms against the 200 ms target**, accounted for in full in § AC4 and § Scope Refinement.

- Frontend: <https://sorobanscan.rumblefish.dev>
- API: <https://api-sorobanscan.rumblefishdev.com/v1>

2. Approved Deliverable

Deliverable 3 - Mainnet Launch:

The explorer is opened to the public on mainnet, load-tested, monitored, and signed off against a security checklist, with a 7-day post-launch monitoring report demonstrating stability.

Acceptance criteria, quoted verbatim from the approved submission:

1. **AC1** — "Block explorer publicly accessible at production URL, showing live mainnet data with ledger sequences matching network tip within 30 seconds"
2. **AC2** — "GitHub repository public; `cdk deploy` from README works in a fresh AWS account"
3. **AC3** — "CloudWatch dashboard accessible to Stellar team (read-only IAM role); all alarms OK; Galexie ingestion lag <30 s from network tip"
4. **AC4** — "Load test report: p95 <200 ms at 1M requests/month equivalent; error rate <0.1%"
5. **AC5** — "Security checklist signed off: no wildcard IAM, RDS has no public endpoint, all secrets in Secrets Manager, all API inputs validated"
6. **AC6** — "7-day post-launch monitoring report: uptime %, API error rate, p95 latency, Galexie ingestion lag per day"

Where the launched system does not match a criterion word for word, the gap is stated in § Scope Refinement rather than reworded here.

3. Live Endpoints and Reviewer Access

Resource	URL / Access
Frontend	https://sorobanscan.rumblefish.dev (public — gate removed)
API base	https://api-sorobanscan.rumblefishdev.com/v1
Swagger UI	https://api-sorobanscan.rumblefishdev.com/api-docs — publicly readable, no key required
OpenAPI JSON	https://api-sorobanscan.rumblefishdev.com/api-docs-json — publicly readable, no key required
API access model	Access-controlled — anonymous requests return 401. The public explorer authenticates transparently at the edge; direct API access is available to reviewers on request (same model as Milestone 2).
Monitoring	CloudWatch dashboard production-soroban-explorer (eu-central-1) with Slack-wired alarms and X-Ray tracing. Read-only access for the Stellar team is available on request.

At launch the frontend is publicly accessible with no Basic Auth gate. The verification video demonstrates the live application end-to-end.

4. Architecture Evidence

The launched system spans three paths, all reproducible from code (AWS CDK + Hetzner Ansible):

1. **Ingestion / write path** (Milestone 1) — Galaxie on ECS Fargate exports mainnet LedgerCloseMeta XDR to S3; a Rust indexer Lambda decodes each object and writes to ClickHouse. End-to-end ledger-close to DB-write is well under the 30-second target.
2. **Read path** (Milestone 2) — Browser → Cloudflare edge → API Gateway → Rust / axum API Lambda → ClickHouse (Hetzner, mTLS). Static SPA: Browser → CloudFront → private S3.
3. **Launch controls** (Milestone 3) — the Cloudflare edge in front of the data API (managed rule sets, rate limiting, bot challenge, DDoS mitigation) plus an origin lock and API Gateway request throttling; CloudWatch dashboard + Slack alarms + X-Ray tracing.

Evidence images:




Figure 1 — Read path as launched: the Cloudflare edge applies managed rule sets, rate limiting and the X-Edge-Secret origin lock (ADR 0048); the API Lambda verifies the Turnstile token and issues the session JWT carried as Authorization: Bearer on every subsequent call.

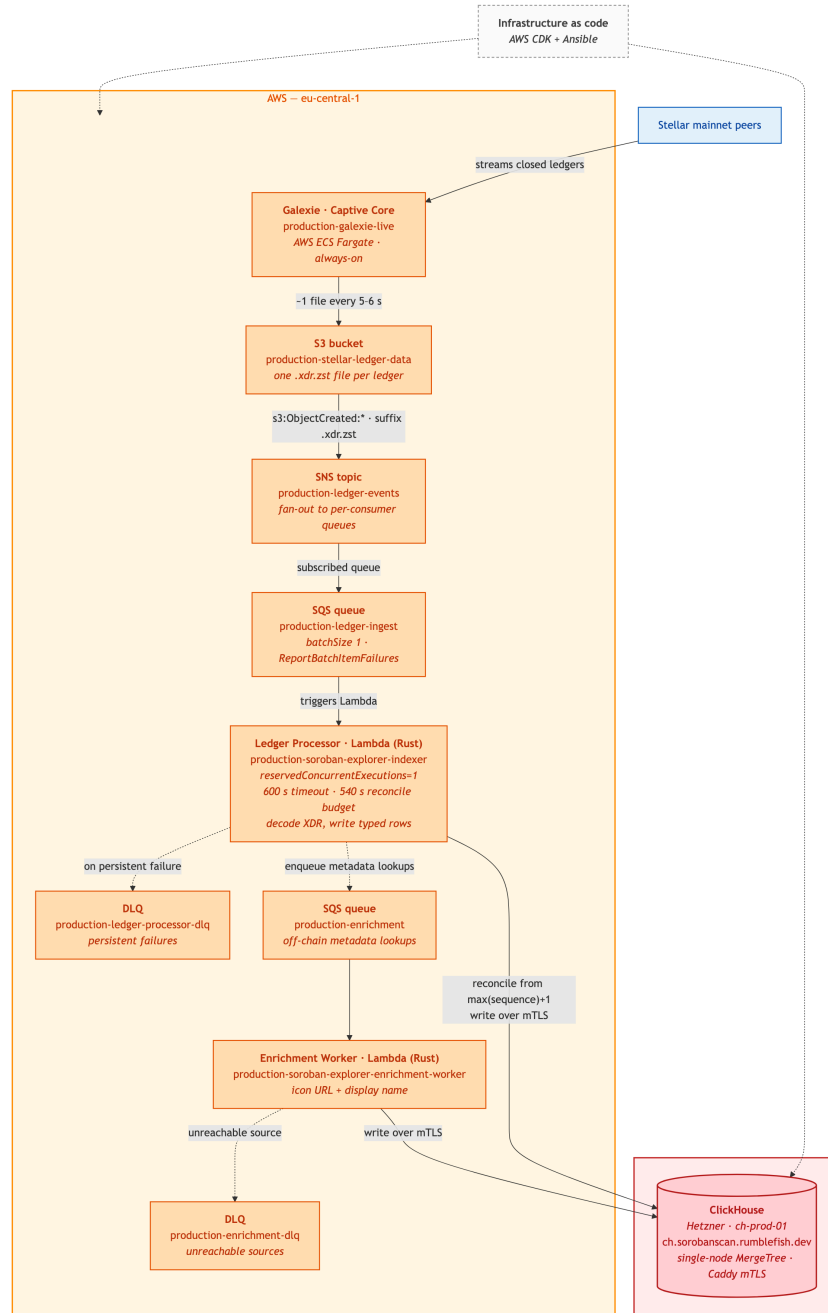


Figure 2 — Ingestion / write path (reused from M1): Stellar mainnet peers → Galaxie on ECS Fargate → S3 → SNS → SQS → Rust indexer Lambda → ClickHouse on Hetzner (mTLS). The SNS topic exists because a second project in the same AWS account subscribes its own queue to the same object-created events. The indexer also enqueues off-chain metadata lookups for a separate enrichment worker; that worker's dead-letter queue is the one alarm raised in § AC3. The CloudWatch dashboard and alarms are shown in § AC3.

5. Acceptance Criteria Evidence

AC1 - Public access, live data ≤30s from tip

Status: met. The pre-launch Basic Auth gate was removed on 2026-07-17 (task 0405). <https://sorobanscan.rumblefish.dev> serves the explorer publicly — HTTP 200 with no WWW-Authenticate challenge.

Data freshness is measured rather than asserted. The indexer publishes IngestionLagSeconds (CloudWatch namespace SorobanBlockExplorer/Indexer, dimension Environment=production): the wall-clock gap between a ledger closing

on mainnet and its row being committed to ClickHouse. Sampled over 75 minutes on 2026-07-17 (15 × 5-minute datapoints):

Measure	Value	Target	Verdict
Ingestion lag, average	3.1 s	< 30 s	Met
Ingestion lag, worst	6.0 s	< 30 s	Met

In steady state the explorer is within ~3 seconds of network tip. Across the full 7-day post-launch window the daily maximum held in the 7–9 s band (§ AC6) — a 3× margin at the worst point measured.

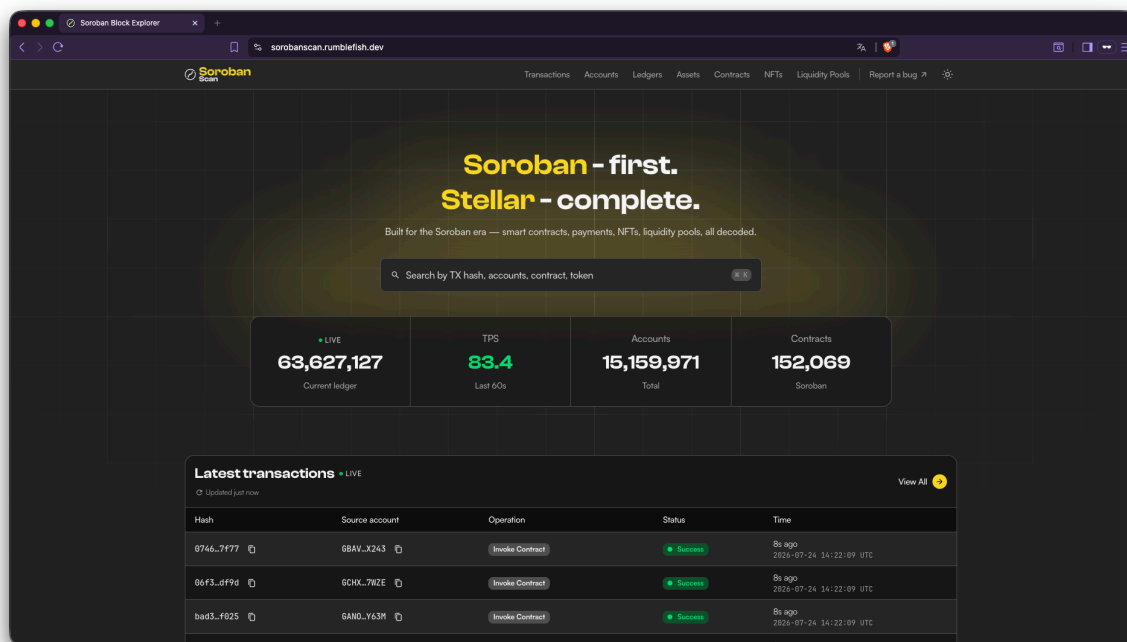


Figure — production URL, no access gate; the newest entries are seconds old, corroborating the measured freshness above.

Data freshness is both **measured** (the IngestionLagSeconds figures above — average 3.1 s, worst 6.0 s) and **visible** in the capture, where the most recent transactions carry timestamps within seconds of real time. The read path is the project's own indexed ClickHouse, not a third-party chain API.

AC2 - Public repo + reproducible cdk deploy

The repository is public and the stack is reproducible from code.

The workspace also carries a substantial automated test suite — **920 test functions** across the Rust crates, including **360 in xdr-parser** (XDR decoding, state-change extraction, and CAP-67 event interpretation — trustline/asset/claimable-balance parsing, mint/burn event handling) and **248 in api** (endpoint responses, OpenAPI schema, pagination). The suite runs with `cargo test`.

- Public repository: <https://github.com/rumblefishdev/soroban-block-explorer>
- Fresh-account deployment runbook: [infra/README.md](#) — prerequisites → make bootstrap → make deploy-production (all AWS stacks) plus the Hetzner Ansible side — with the operational guide in [docs/deployment.md](#). A few steps are manual by design: ordering the Hetzner dedi-

cated server and Storage Box, and publishing the box IPv4 to SSM (/soroban/production/ch-ip).

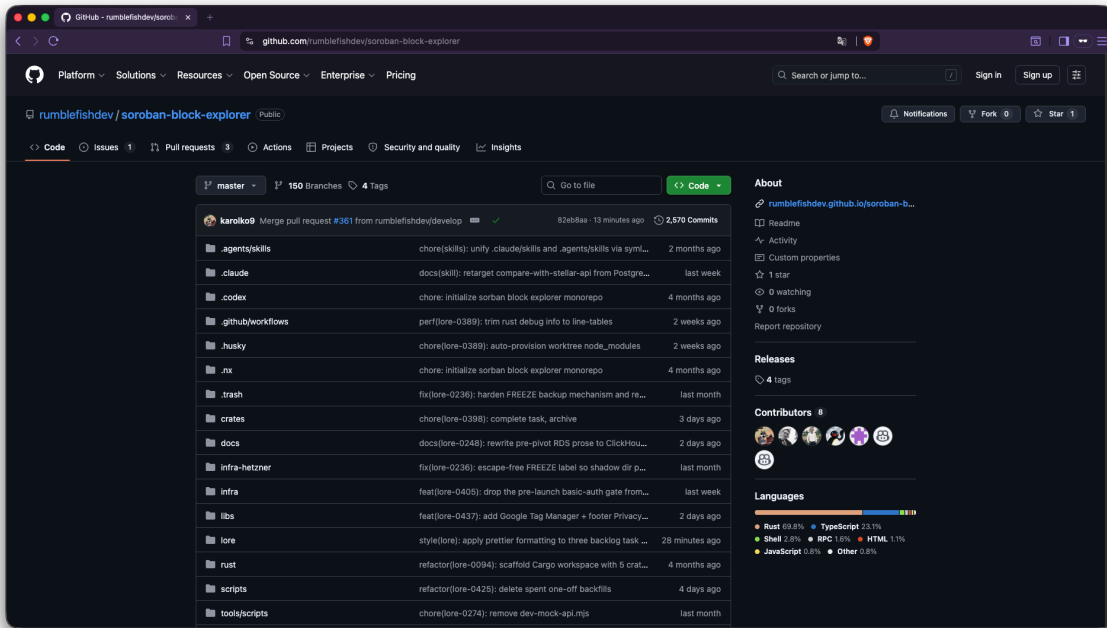


Figure — `github.com/rumblefishdev/soroban-block-explorer` is public and carries the Rust / ClickHouse / infrastructure code.

AC3 - Monitoring dashboard, ingestion lag

CloudWatch dashboard with Slack-wired alarms and X-Ray is live.

Ingestion lag under 30 s: met and measured. The indexer emits SorobanBlockExplorer/Indexer / IngestionLagSeconds per ledger written (task 0399, live in production) — average 3.1 s, worst 6.0 s against the < 30 s criterion (figures and method in AC1). This metric is also the per-day source for the AC6 report.

The dashboard's **API** section graphs API Lambda latency (p50/p95/p99) and API Gateway 4xx/5xx error counts alongside the ingestion panels; the Cloudflare edge and API Gateway throttling protect the public ingress (see AC5). The alarms on the tracked metrics — API 5XX-rate, Galexie ingestion-lag, ledger-processor error-rate, and ClickHouse-write failures — are in OK state.



Figure — the production dashboard: ingestion and API health on a single pane.

CloudWatch Alarms (28)						
Name	State	Actions	Actions muted - new	Last state update (UTC)	Conditions	
production-ledger-processor-error-rate	OK	Actions enabled	-	2026-05-27 11:45:38	Error Rate > 0.01 for 1 datapoints within 5 minutes	
production-ledger-processor-dlq-depth	OK	Actions enabled	-	2026-07-24 14:50:55	DLQ depth > 0 for 1 datapoints within 1 minute	
production-indexer-ch-write-failures	OK	Actions enabled	-	2026-05-27 11:41:20	ClickHouseFailures > 10 for 1 datapoints within 5 minutes	
production-galaxie-ingestion-lag	OK	Actions enabled	-	2026-07-13 09:40:35	NumberOfMessagesSent < 1 for 1 datapoints within 5 minutes	
production-galaxie-sphereml-storage	OK	Actions enabled	-	2026-07-02 13:49:26	Ephemeral Used (Mi) > 60 for 3 datapoints within 15 minutes	
production-enrichment-worker-error-rate	OK	Actions enabled	-	2026-07-15 12:39:42	Worker Error Rate > 0.01 for 1 datapoints within 5 minutes	
production-enrichment-dlq-depth	In alarm	Actions enabled	-	2026-07-03 09:02:06	Enrichment DLQ depth > 0 for 1 datapoints within 1 minute	
production-api-gateway-5xx-rate	OK	Actions enabled	-	2026-07-16 19:09:44	5xx Rate (%) > 0.5 for 1 datapoints within 5 minutes	
prices-production-supply-errors	OK	Actions enabled	-	2026-07-21 13:16:32	Errors >= 1 for 1 datapoints within 1 hour	
prices-production-index-push-freshness	OK	Actions enabled	-	2026-07-08 12:15:18	PushAggSeconds > 604800 for 1 datapoints within 15 minutes	
prices-production-oracle-errors	OK	Actions enabled	-	2026-07-23 14:58:56	Errors >= 1 for 1 datapoints within 5 minutes	
prices-production-mtms-notifier-probe-no-innovations	OK	Actions enabled	-	2026-07-21 13:01:46	Innovations < 1 for 3 datapoints within 3 days	
prices-production-mtms-notifier-probe-errors	OK	Actions enabled	-	2026-07-06 10:17:42	Errors >= 1 for 1 datapoints within 1 day	
prices-production-mtms-notifier-probe-duration-near-timeout	OK	Actions enabled	-	2026-07-21 13:18:54	Duration >= 48000 for 2 datapoints within 2 days	
prices-production-mtms-notifier	OK	Actions enabled	-	2026-07-08 12:30:16	MinDaysToNextRun < 30 for 1 datapoints within 1 day	
prices-production-ledger-processor-no-innovations	OK	Actions enabled	-	2026-07-09 08:35:18	Innovations < 1 for 1 datapoints within 15 minutes	
prices-production-ledger-processor-lag	OK	Actions enabled	-	2026-07-15 22:05:45	ApproximateAgeOfOldestMessage > 120 for 5 datapoints within 5 minutes	
prices-production-ledger-processor-errors	OK	Actions enabled	-	2026-07-15 09:00:00	Errors >= 1 for 1 datapoints within 5 minutes	
prices-production-ledger-processor-dlq	OK	Actions enabled	-	2026-07-08 08:07:19	ApproximateNumberOfMessagesVisible >= 1 for 1 datapoints within 5 minutes	
prices-production-enrichment-no-innovations	OK	Actions enabled	-	2026-07-21 12:36:54	Innovations < 1 for 3 datapoints within 3 hours	
prices-production-enrichment-errors	OK	Actions enabled	-	2026-07-21 12:40:22	Errors >= 1 for 1 datapoints within 1 hour	
prices-production-enrichment-duration-near-timeout	OK	Actions enabled	-	2026-07-24 11:06:17	Duration >= 240000 for 2 datapoints within 2 hours	

Figure — production alarms, captured 2026-07-25. The API 5XX-rate, Galaxie ingestion-lag, ledger-processor error-rate and ClickHouse-write alarms are in OK state. prices- entries belong to a separate project in the same AWS account. One alarm is raised — production-enrichment-dlq-depth, which fires whenever the metadata-enrichment queue is non-empty. That queue retries two cosmetic fields, an icon URL and a display name, against sources the asset issuer controls; 97 % of assets carry them and the rest have no reachable source. No chain data is involved, and no metric reported in § AC6 depends on it.

AC4 - Load-test report

Status: partially met — error rate met with a ~100× margin; p95 not met. This section states the result plainly and then accounts for it in full.

AC4 half	Target	Measured (1M-req/ month equivalent)	Verdict
Error rate	< 0.1 %	0.000 % (0 failures)	Met
p95 latency	< 200 ms	577 ms	Not met

Median latency (p50) is **168 ms** at the required load and **151 ms** at 40× that load — the typical request is inside the target. The miss is in the tail.

Method: why the load is expressed as a rate, not as virtual users

The criterion is "1M requests/month equivalent" — a **request rate** (1,000,000 / 30 days ≈ **0.4 req/s**). The harness therefore drives an **open model**: requests arrive at a fixed rate (Poisson), independent of how fast the server answers.

The earlier plan called for "≥1000 concurrent users". That framing was abandoned for a measured reason, not for convenience: a virtual-user driver is **closed-loop**, so its request rate is an *output*, not an input — each user waits for a response before sending the next. Our own measurement of the same --vus 4 configuration produced **9,459 req/s against a local stub** and **~10 req/s against production**. A VU count therefore cannot express "1M requests/month"; only an arrival rate can. Open-model testing is what actually measures the stated criterion. (Recorded in task 0357.)

Result across load tiers

Three open-model runs were made against the production API (**16,232 requests** in total). The runs targeted the API Gateway origin rather than the Cloudflare-fronted hostname, so the figures measure the backend itself without the edge's own rate limiting; the 40× tier additionally ran with the per-IP AWS WAF rate rule lifted, since that rule exists precisely to stop this traffic pattern from a single source. (That WAF has since been retired — ADR 0048.) These map to the proposal's two documented load points — the **1.2M/month tier is the "1M baseline"** and the **10.2M/month tier is the "10M stress"** run — plus an additional 40× (**49.3M/month**) capacity check. Representative tiers:

Load tier	Rate	Requests	Error rate	p50	p95
1.2M req/ month (AC4 level)	0.48 /s	168	0.000 %	168 ms	577 ms
10.2M req/ month (8×)	3.95 /s	2,363	0.000 %	149 ms	567 ms
49.3M req/ month (40×)	19.04 /s	13,701	0.000 %	151 ms	575 ms

The p95 is flat across a 40× range of load (577 ms → 575 ms). This is the central finding: the p95 is **not load-dependent** in the tested range. The system absorbs forty times the required traffic with an unchanged tail and zero failures — capacity is proven far beyond the criterion. The p95 miss is therefore **not a capacity problem**; it is a set of fixed per-request costs that no amount of headroom removes.

Where the 577 ms comes from

Decomposing the same measurement by endpoint class:

Scope	p95 @ AC4 load	p95 @ 40× load
All 26 endpoints	577 ms	575 ms
Excluding the 2 synchronous-external-fetch endpoints	342 ms	506 ms
Excluding those and lplist	280 ms	317 ms

Three distinct causes, in order of size:

1. **Synchronous external fetches on two detail endpoints (largest).** txdetail and nftdetail fetch data from outside our infrastructure while the request is open. txdetail spends **53 ms in ClickHouse and 1,074 ms (p95) waiting on an S3 archive read in us-east-2 from a Lambda in eu-central-1** — a cross-region round trip per request. nftdetail spends **96 ms in ClickHouse and 1,604 ms (p95)** on a third-party IPFS gateway plus a token_uri() RPC. These are **deliberate architectural decisions** (ADR 0029, ADR 0043: heavy/detail-only fields are fetched at read time rather than persisted), not defects — but they place the tail of those two endpoints partly outside our control.
2. **A known ClickHouse cost on one endpoint.** lplist derives each listed pool's creation ledger by scanning that pool's full snapshot history — the endpoint reads **11.6M rows per request** (mean of per-request read_rows for lplist at the 40× tier, results.csv). This is our own technical debt: the design decision that removed the stored column (task 0208) assumed this read would be cheap; measurement shows it is not. The cause is identified to the query; the candidate fix is tracked in task 0401.
3. **A fixed infrastructure floor (irreducible today).** Every request pays ≈ 40 ms before any data is read (API Gateway + Lambda + mTLS to ClickHouse), plus ≈ 15 ms per ClickHouse round trip. Measured across all 26 endpoints, the median non-database overhead fits ≈ 40 ms + 15 ms $\times$ (number of CH queries) — from 51 ms for a single-query endpoint to 157 ms for an 8-query endpoint. An endpoint issuing five queries has **already spent ~ 115 ms of the 200 ms budget** before ClickHouse returns a single row.

Why pay for those round trips at all?

Because the machine that serves them costs 9.5× more inside AWS. ClickHouse runs on a dedicated server — 24 threads, 128 GB RAM, 1.72 TiB of local NVMe holding our full 929 GiB working set, plus off-box backups — for **\$133/month, AWS egress included**. The closest on-demand configuration in eu-central-1 — 128 GiB of RAM and ≈ 1.9 TB of local NVMe — is an m6gd.8xlarge at **\$1,273/month**, backups extra. The round-trip cost in cause 3 is the bill for that: at a mean of 3.29 ClickHouse queries, ≈ 50 ms per request. (AWS list prices, 2026-07-27.)

It is also not why AC4 misses. Remove every round trip outright — a strictly generous assumption about what moving into the VPC would recover — and p95 moves 577 $\rightarrow$ 547 ms: **8 % of the 377 ms gap**. The rest is the three causes above.

Optimisation performed during this milestone

The load test drove a measured optimisation pass on the read path. Total ClickHouse work per test run fell from **78.3 billion to 23.89 billion rows (–69 %)** at identical load, via two query fixes ([#347](#), [#349](#)) and one index. Endpoints that were **not modified at all** improved by 3–4× as a direct result (ldgdetail p95 890 $\rightarrow$ 319 ms, ldglist 1015 $\rightarrow$ 229 ms) — evidence that the earlier latency was contention from a saturated database rather than the endpoints themselves.

Raw measurement data: [load-tests/](#) holds the three tiers quoted above — per tier a results.csv (one row per request: client latency joined to the ClickHouse query log for that same request via request_id, so each request carries its own read_rows / ch_queries / ch_duration_ms) and a client_summary.csv (per-endpoint percentiles). The p50 / p95 figures in this section are percentiles of duration_ms across a tier's results.csv; [load-tests/README.md](#) maps each tier to its run and restates the headline numbers.

Honest assessment

Restated plainly for the reviewer: **the p95 < 200 ms criterion is not met at any tested tier, including the required one — 577 ms against a 200 ms target**. It is not met even with all three causes above excluded (280 ms). What the measurement does establish:

- the **error-rate half of AC4 is met with a $\sim 100\times$ margin** (0.000 % vs 0.1 %);
- the **median request meets the latency target** (168 ms at the required load);
- **capacity is proven at 40× the required load** with no degradation and no errors;

- every millisecond of the gap is **attributed to a named, measured cause**; follow-ups are tracked in tasks 0401 (lplist) and 0402 (per-query overhead).

AC5 - Security checklist signed off

All ingress and application controls are in place and verified: least-privilege IAM (no wildcard actions), the Cloudflare edge in front of the data API with an origin lock and API Gateway request throttling, a Turnstile-to-session-JWT gate on the data endpoints, no public datastore endpoint (ClickHouse published to loopback only — 127.0.0.1:8123/:9000 — behind Caddy mTLS on a firewalled host that admits only ports 22/80/443), secrets in AWS Secrets Manager, TLS end-to-end, server-side input validation on every endpoint, SSE-S3 at rest on the public ledger bucket, and automated weekly off-box ClickHouse backups.

The full control list and an OWASP Top 10 (2021) coverage mapping are in [milestone-3-security-checklist.md](#). The original KMS-at-rest and point-in-time-recovery items were RDS-specific and are met by equivalents after the RDS retirement — see § Scope Refinement.

```
<E "ssh://[^\''"]+ /usr/local/bin/ch-backup | head -1)"
#Ansible: ch-backup
30 3 * * 0 root /usr/local/bin/ch-backup >> /var/log/ch-backup.log 2>&1

ch-20260705T035828Z      Sun, 2026-07-05 03:58:33 [d768f611c7bdcc883fc19f8c65a72ae305a1b00f9fe14de85b5663f572096e2f]
ch-20260712T035324Z      Sun, 2026-07-12 03:53:30 [3f50f095db810c56f67937f85d1184f72a6603cf04f67709c0da5ec290a253e2]
ch-20260719T035534Z      Sun, 2026-07-19 03:55:40 [4ff6b08d64d8f81c144b8207da16cbf4d7ca337c693ec0db4d84da0c4e85a55d]
ch-20260726T033815Z      Sun, 2026-07-26 03:38:22 [ed0f2930c3ff0343de8f7cd3ffd9019f1c3306a58e71426adcfa398349471626]
```

Figure — backups, verified on the host on 2026-07-27: a Sunday cron entry and four archives seven days apart, evidencing the weekly cadence and the four-week retention. The full run log is quoted in the security checklist.

AC6 - 7-day post-launch monitoring report

Generated from production telemetry over the first 7 days after launch.

- Full report, per-day figures and the metric queries behind them: [milestone-3-7day-report.md](#)

The window ran **2026-07-17 13:40Z → 2026-07-24 13:40Z**. Per-day figures are in the report; the window summary:

Metric	Result over the window	Target	Verdict
Uptime	100.00 % (derived — see report)	≥ 99.9 %	Met
API error rate	0.000 % — zero 5XX out of 30,378 requests served	< 0.1 %	Met
Ingestion lag	7–9 s (worst 9 s)	< 30 s	Met
Ledger completeness	0 gaps, all 7 days	0 gaps	Met
API p95 latency	553 ms worst day (4 of 7 days < 200 ms)	< 200 ms	Not met

The p95 line is the same criterion as AC4 and misses for the same reasons. Note this report measures API Gateway Latency (server-side), which excludes the client network leg the load test includes — hence the lower figures here, with four of seven days under target. Uptime is derived from zero 5XX responses and from production-api-gateway-5xx-rate holding OK throughout; the report records which alarms were raised over the window and why they do not bear on these metrics.

Scope Refinement — Deviations from the Approved Plan

Three points where the launched system differs from the original Deliverable 3 wording, stated openly (following the Milestone 1 precedent, which recorded the PostgreSQL-on-RDS → ClickHouse-on-Hetzner change the same way):

1. **p95 latency target not met (AC4).** Measured p95 is 577 ms against the 200 ms target. It is not a capacity limit — the tail is flat across a 40× load range with zero errors — but a set of fixed per-request costs, each measured and attributed in § AC4. The error-rate half of AC4 is met with a ~100× margin.
2. **Data-at-rest / recovery wording was RDS-specific; RDS is retired.** The original security checklist named KMS-at-rest and point-in-time recovery, both specific to the PostgreSQL-on-RDS datastore that was retired (task 0239) in favour of ClickHouse on Hetzner. They are satisfied here by architecture-appropriate equivalents: the public ledger bucket is encrypted with SSE-S3 (AES256 — deliberately not SSE-KMS, since the contents are public on-chain XDR), and the ClickHouse store — which holds only public, fully re-derivable chain data — is protected by automated weekly off-box backups. No datastore is publicly reachable (ClickHouse is bound to loopback behind mTLS/Caddy on a firewalled host).
3. **AC3 — one alarm raised.** The criterion asks for "all alarms OK". The four alarms on the criteria-bearing metrics — API 5XX rate, Galexie ingestion lag, ledger-processor error rate, ClickHouse write failures — are OK. One further alarm, production-enrichment-dlq-depth, is raised: it fires whenever the metadata-enrichment queue is non-empty, and bears on no AC6 metric — ledger completeness is independently verified at 0 gaps on every day.

6. Source References

Resource	Link
Public repository	rumblefishdev/soroban-block-explorer
Technical design	technical-design-general-overview.md
Infra / deploy	infra/README.md
Load-test harness	crates/load-tests
M1 / M2 evidence	milestone-1-evidence.md · milestone-2-evidence.md